

Lecture-10-11

Course: Data Science



Introduction to Neural Network and Deep Learning

By

Dr. Sibarama Panigrahi

Senior Member, IEEE

Assistant Professor, Department of Computer Sc. & Engineering
National Institute of Technology, Rourkela, Odisha, 769008, India

Mobile No.: +91-7377302566

Email: panigrahis[at]nitrkl[dot]ac[dot]in
panigrahi[dot]sibarama[at]gmail[dot]com


Outlines

- Introduction to Neural Network
- **Mathematical Model for Neural Network**
- Differentiation and its Application to Train Neural Network
- **Deep Neural Network**
- Recent Advances in Deep Learning
 - Activation Function, Weight Initialization (Xavier & Glorot, He)
 - Dropout and Regularization, Batch Normalization
 - Optimizers (SGD, NAG, AdaGrad, AdaDelta, RMSPROP, ADAM)
 - Building and Training Deep Neural Network using Python
- **Introduction to Hyper-Parameter Optimization**

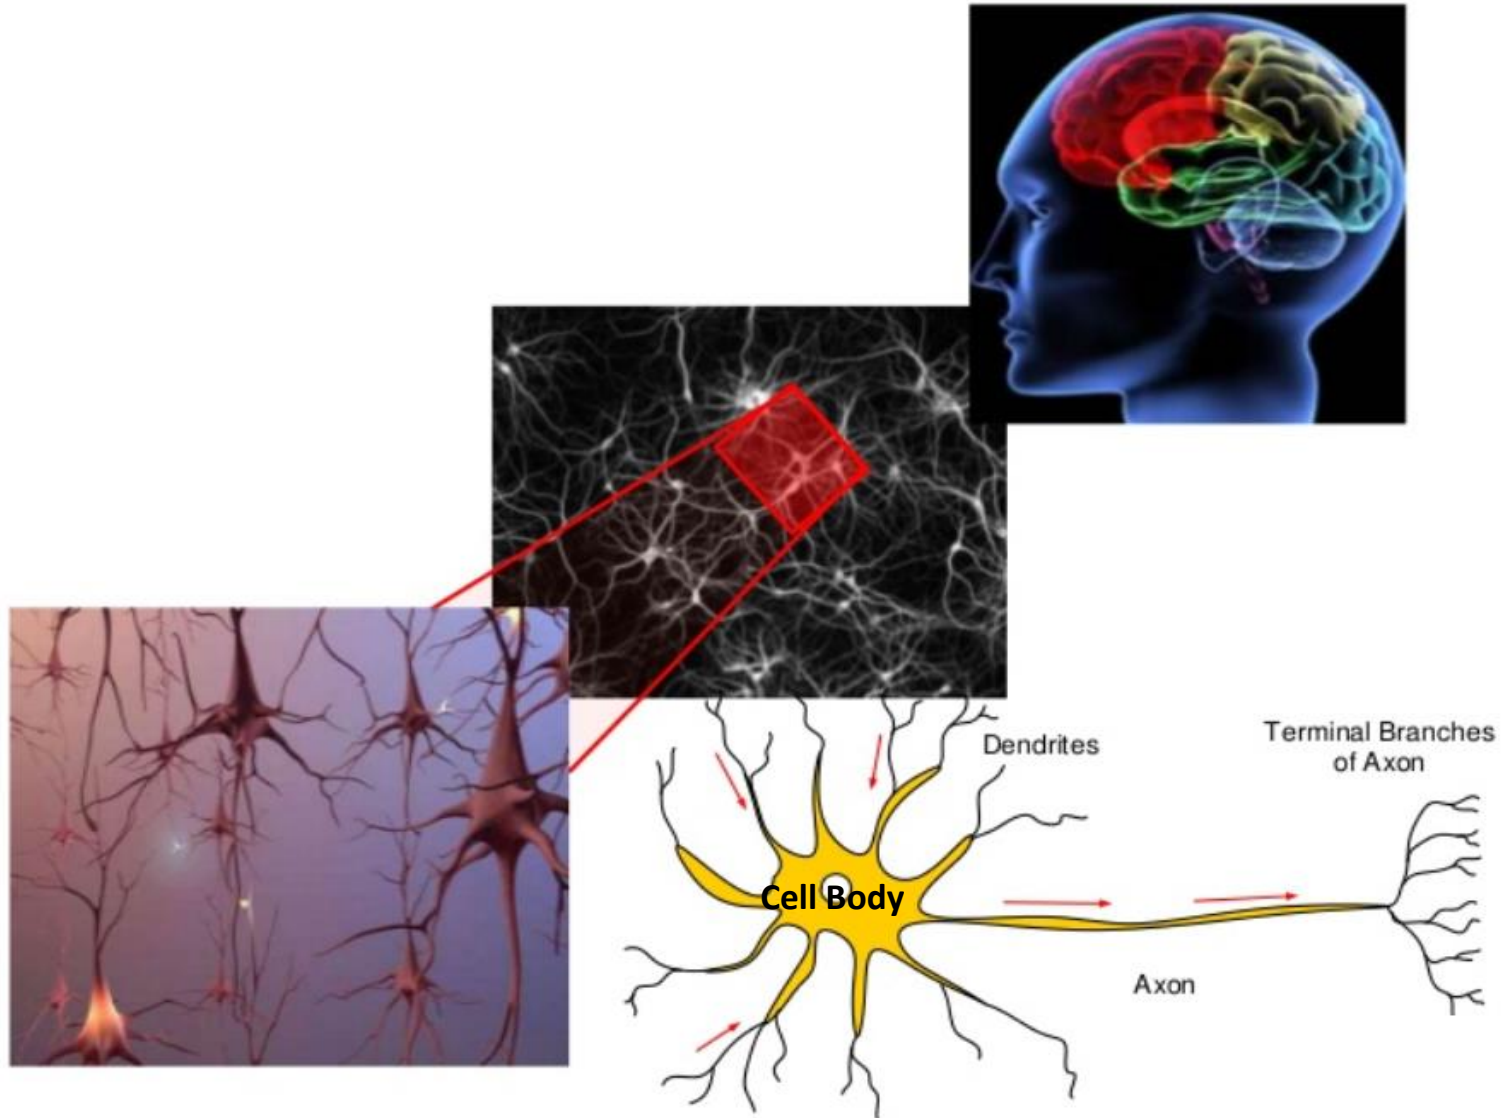
Artificial Neural Network

- An Artificial Neural Network (ANN) is a mathematical model that *loosely simulates* the structure and functionality of **Biological** nervous system to map the inputs to outputs.


Block Diagram of Biological Nervous System




Typical Human Brain



Human Brain Neuron vs Artificial Neuron




Artificial Neuron



$$V_k = W_{k1} * x_1 + W_{k2} * x_2 + W_{k3} * x_3 + \dots + W_{kn} * x_n + b_k$$


Artificial Neuron



$$V_k = \sum_{j=1}^n (W_{kj} * x_j) + b_k$$

$$y_k = f(V_k)$$

Single Neuron Model



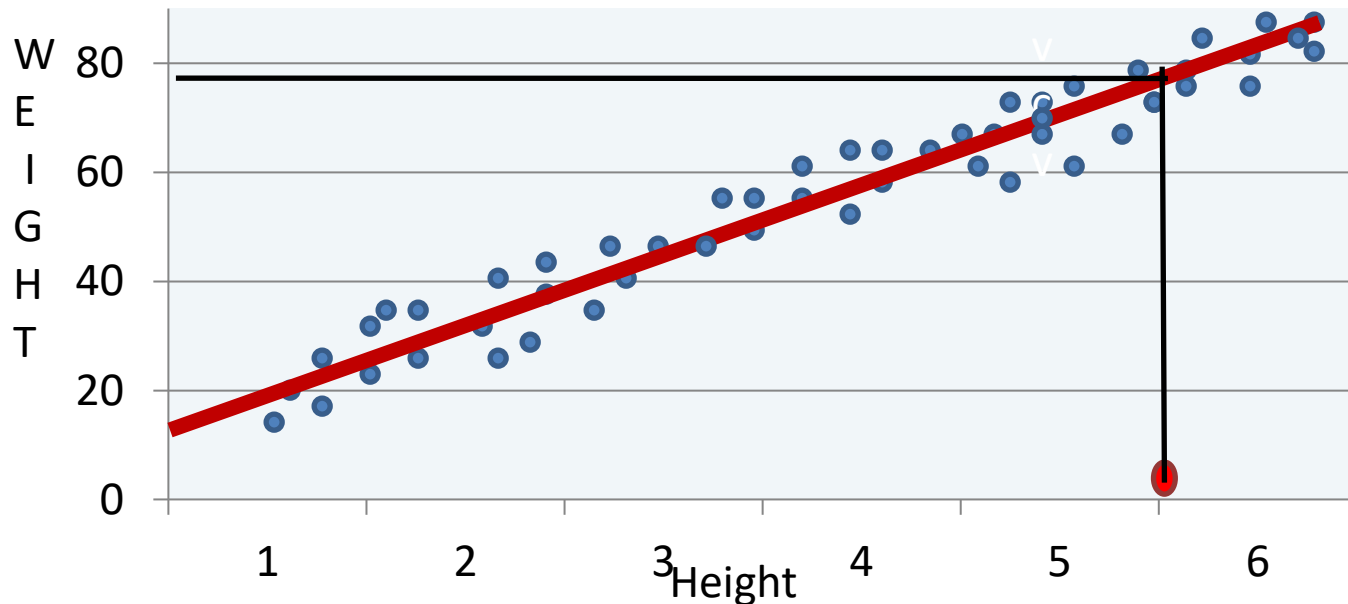
Single Neuron Model

- Application

$$y = mx + c$$

Where m = Slope Of Straight Line


X = Height c = Intercept y = Weight



Single Neuron Model

Error Calculation


- The error $E_i = (\text{Actual Value} - \text{Predicted value}) = (T_i - y_i)$
- For making +ve $E_i = (T_i - y_i)^2$ [Error for i^{th} input instance]



Linear Neural Network

- **Error Calculation**


- It is done to adjust the slope(m) and intercept for better fitting next time.




Linear Neural Network

$$y_k = w_{k1} * x_1 + b_k$$

$$y = m * x + c$$




Output is
Linearly
Dependent on
Input
Parameters



$$v_k = w_{k1} * x_1 + b_k$$

$$y_k = f(v_k) = w_{k1} * x_1 + b_k$$


Plotting Error



Differentiation...

$$y = f(x)$$

$$\frac{dy}{dx} = \frac{df}{dx} = y' = f'$$



How much does y change as x changes = $\frac{\Delta y}{\Delta x} = \frac{y_2 - y_1}{x_2 - x_1} = \frac{p}{b} = \tan(\theta)$

$$\frac{dy}{dx} = \lim_{\Delta x \rightarrow 0} \frac{\Delta y}{\Delta x}$$

Differentiation...

$$\frac{dy}{dx} = \lim_{\Delta x \rightarrow 0} \frac{\Delta y}{\Delta x}$$

As $\Delta x \rightarrow 0$ we obtain a tangent at x .





$$\frac{dy}{dx} = \tan(\theta) = \text{slope of the tangent at } x = x_1$$

$$\frac{dy}{dx} = \text{Slope of the tangent to x-axis at } x = x_1$$

Differentiation...



$$0 < \theta < 90 \quad \tan(\theta) = +ve$$

$$90 \quad \tan(90) = \text{Undefined}$$




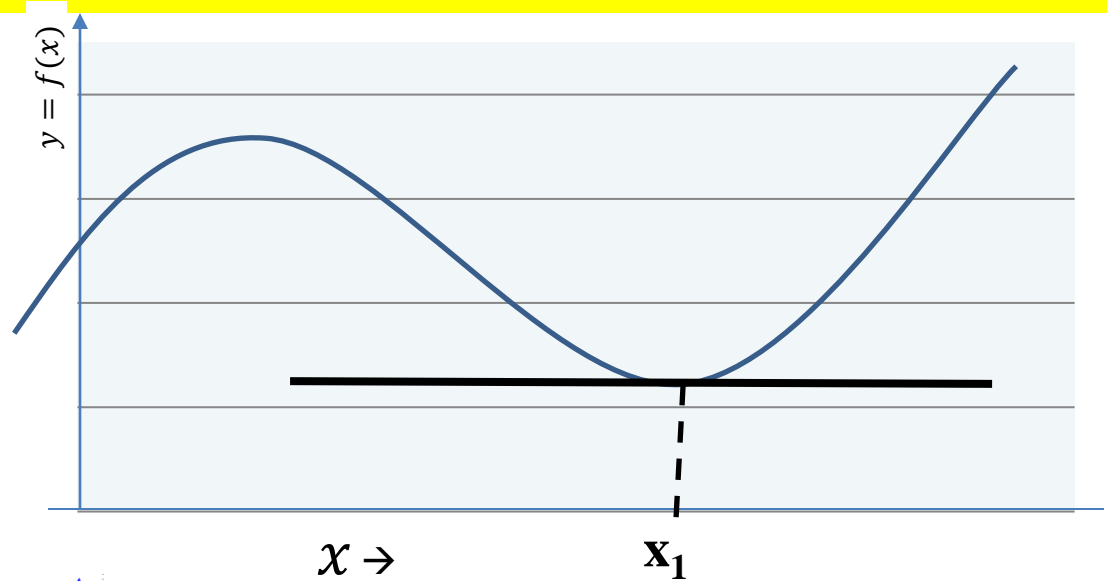
Differentiation...

$$\theta > 90 \quad \tan(\theta) = -ve$$



Differentiation...


$$\theta = 0 \quad \tan(\theta) = 0$$



Differentiation...


$$\theta = 0 \quad \tan(\theta) = 0$$

Maxima



Minima

Note: At minima and Maxima the Slope is 0 $\Rightarrow \tan(\theta) = 0 \Rightarrow \frac{dy}{dx} = 0$



Differentiation...

Distinguishing between a Minima & Maxima

$$\text{Let } f(x) = X^2 - 3X + 2$$

$$\frac{df}{dx} = 0$$

$$2X - 3 = 0$$

$$X = 1.5$$


$$f(1.5) = -0.25$$

Take a point near 1.5, let $X=1$


$$f(1) = 1 - 3 + 2 = 0$$

$X=1.5$ can't be maxima. It is a minima.


Error Function with Minima and No Maxima



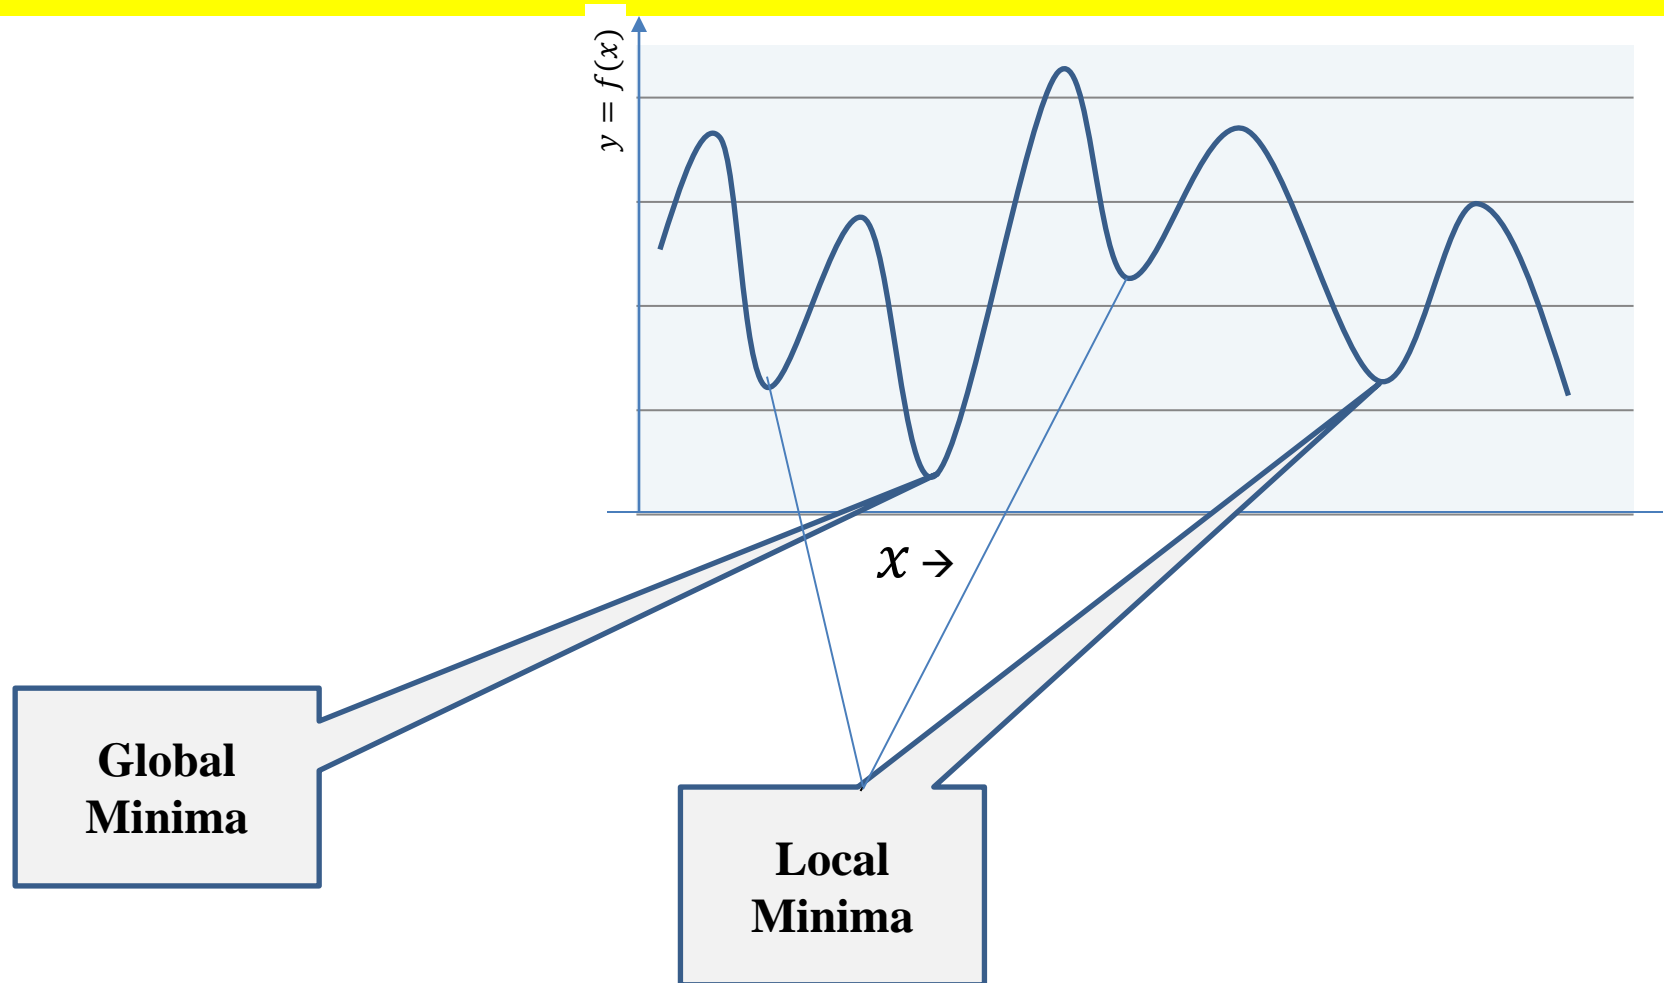
Error Function with a Maxima and No Minima



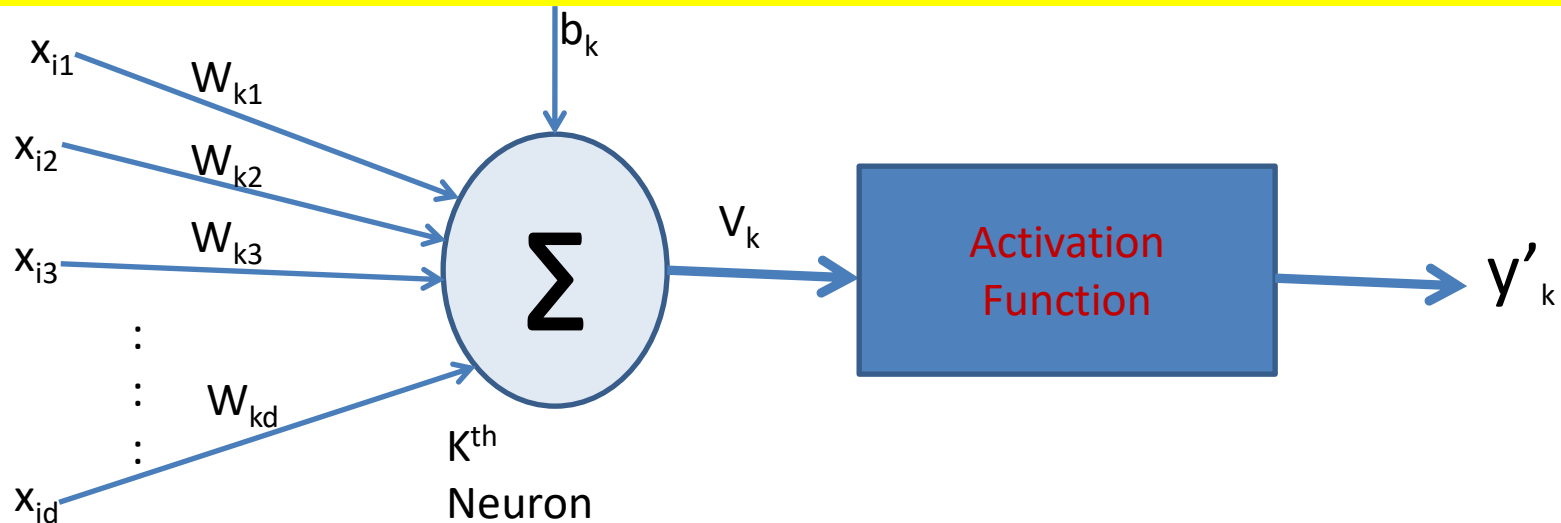
Error Function without a Maxima and Minima



Error Function with multiple Maxima and Minima



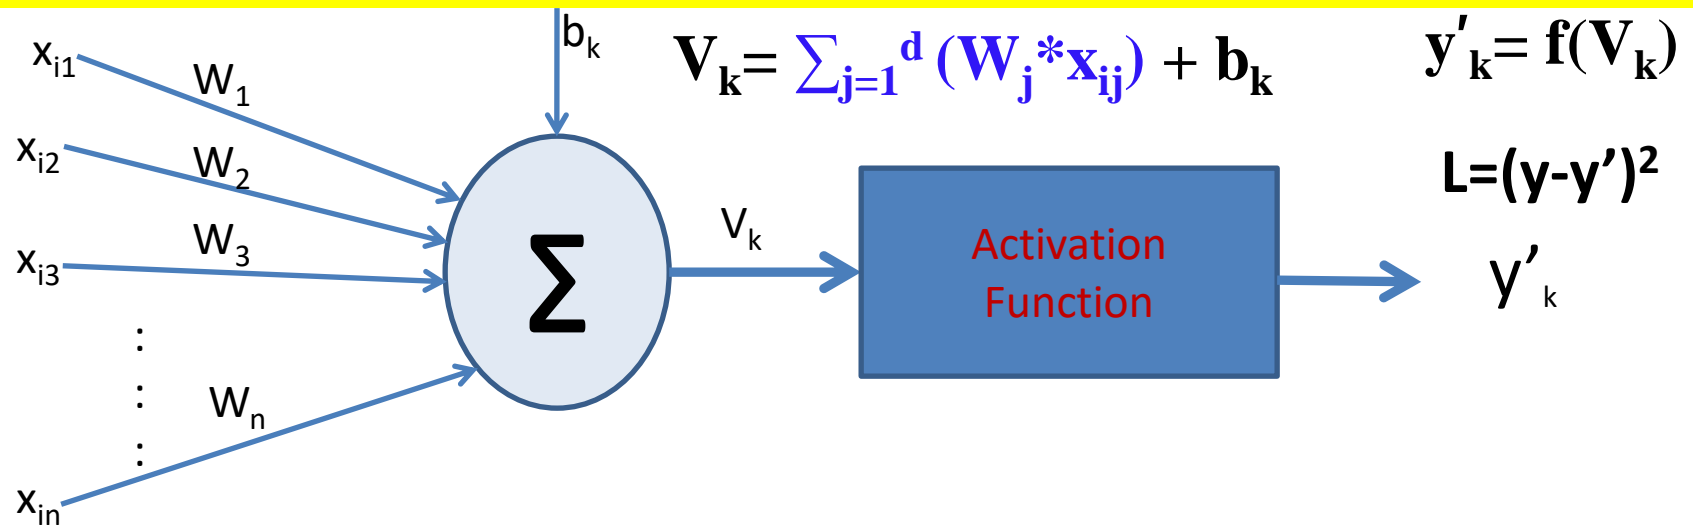
TRAINING A SINGLE-NEURON MODEL



$$V_k = \sum_{j=1}^d (W_{kj} * x_{ij}) + b_k \quad y'_k = f(V_k) \quad L = \sum_{i=1}^n (y_i - f(w^T x_i + b))^2$$

- Step-1: Define the loss function $\sum_{i=1}^n (y_i - y'_i)^2$
- Step-2: Define the optimization $\underset{\tilde{w}_i}{\text{Min}} \sum_{i=1}^n (y_i - f(w^T x_i + b))^2 + \text{reg}$

TRAINING A SINGLE-NEURON MODEL



• Step-3: Solve the optimization problem

- Randomly initialize the weights
- Feed forward the inputs and compute the loss function
- Update the weights

$$(w_i)_{new} = (w_i)_{old} - \eta \left[\frac{dL}{dw_i} \right] = \begin{cases} \frac{\partial L}{\partial w_1} \\ \frac{\partial L}{\partial w_2} \\ \frac{\partial L}{\partial w_d} \end{cases}$$

$$\frac{dL}{dw_1} = \boxed{-2(y - y') * x}$$

$$\boxed{-2(y - y')}$$

$$\boxed{x}$$

TYPES OF NEURAL NETWORK




Figure 1.2: Single layer Neural Network





Figure 1.3: Multilayer Neural Network

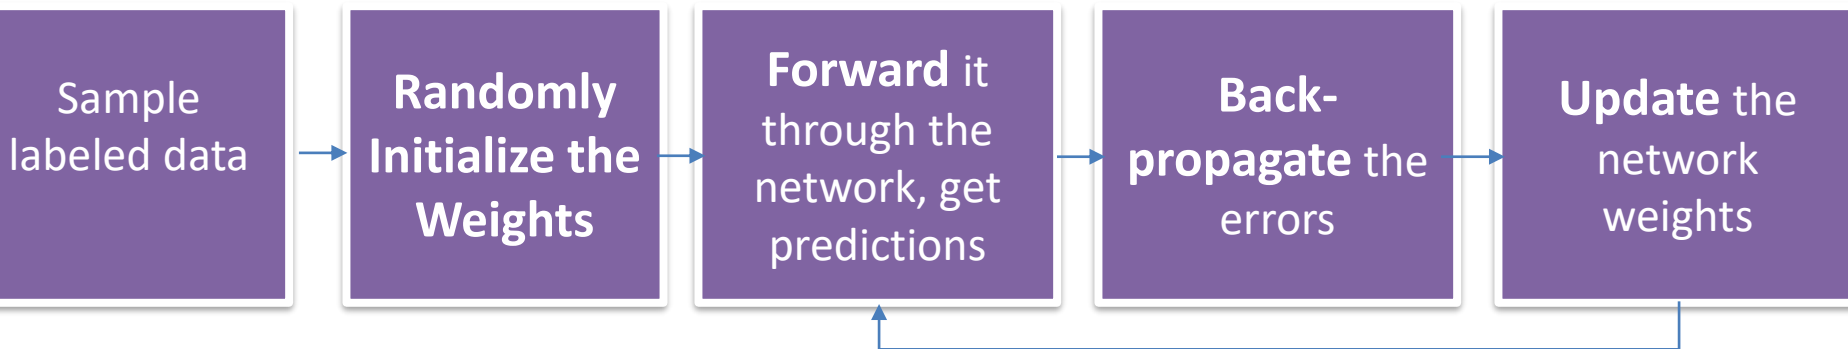
WHY MULTILAYER NEURAL NETWORK?

- **Biological Inspiration**
- **Universal Approximators:** Can approximate any nonlinear function to any desired level of accuracy.
- **Results in Powerful Models**

Graph for $2\sin(x^2) + \sqrt{x^5}$

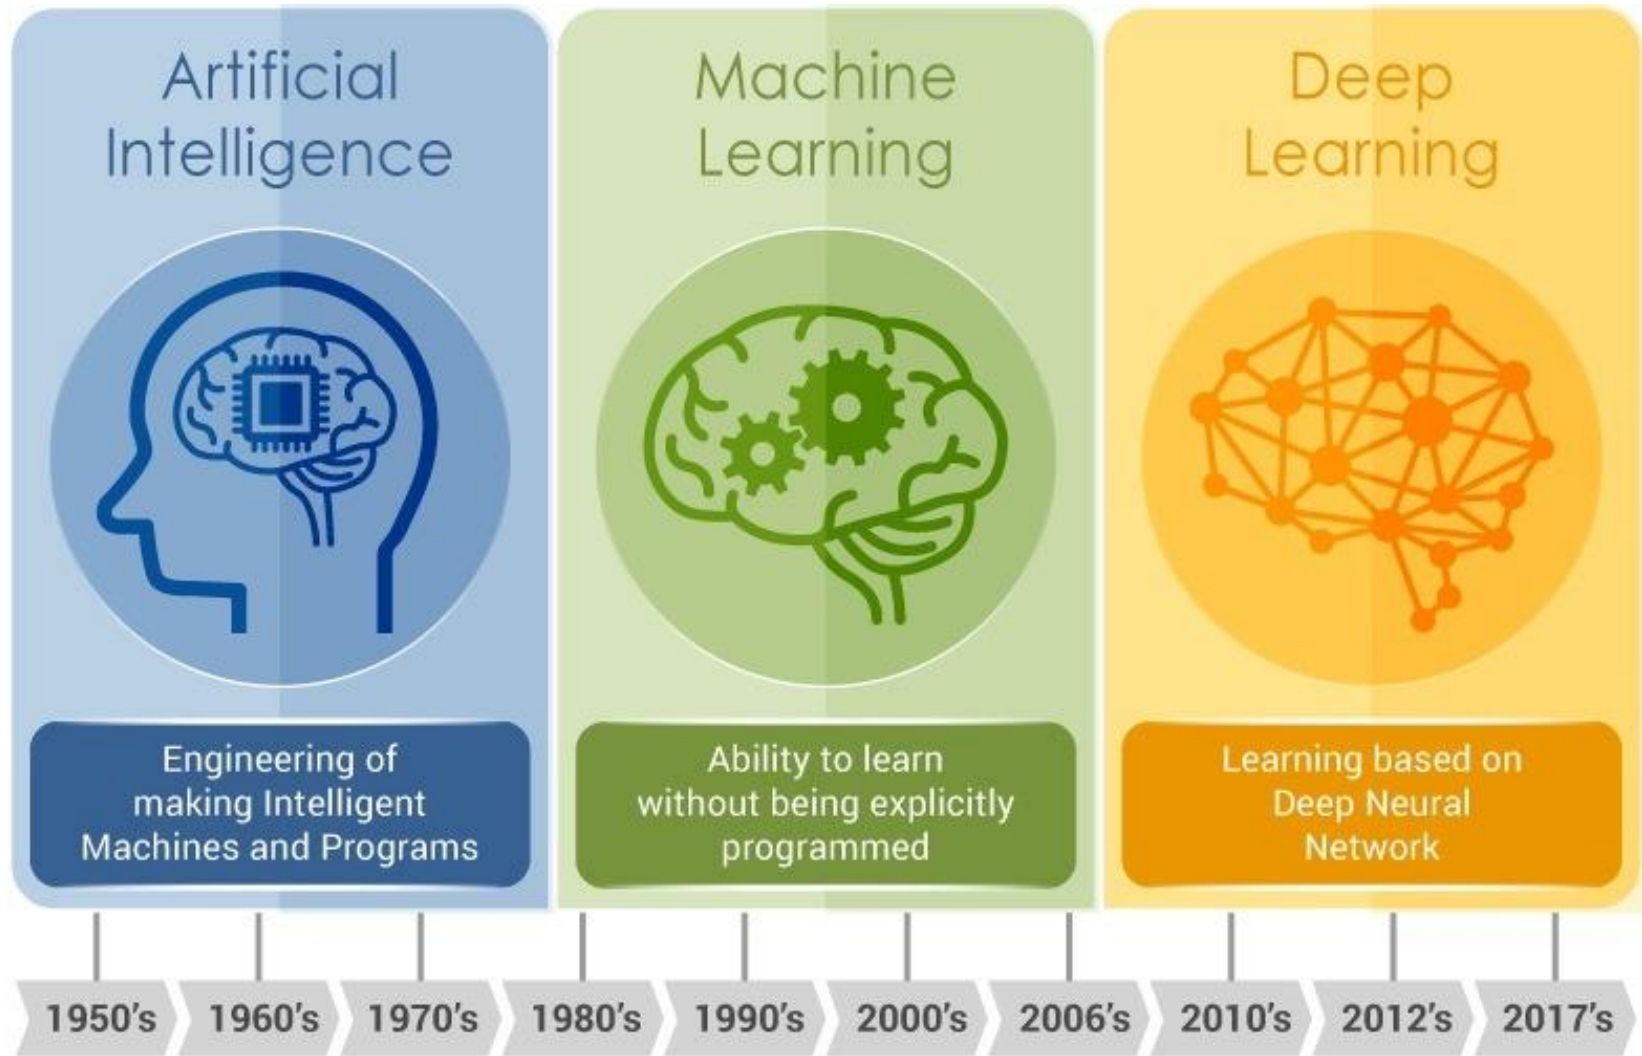


TRAINING MULTILAYER NEURAL NETWORK



- **Back-Propagation: Chain Rule + Memoization**
 - In Stochastic Gradient Descent (SGD) U take one point (Input Vector)
 - In Mini-Batch SGD, U take a set of points(input vectors)
 - In Gradient Descent, U take all the input vectors

AI vs Machine Learning vs Deep Learning



Deep Learning

- A type of *machine learning* based on *artificial neural networks* in which *multiple layers of processing* are used to *extract progressively higher level features* from data.

- “Deep Learning with Python” Francois Chollet



DEEP LEARNING APPROACH

- Standard Approach (Mathematicians)
 - Build new theories
 - Perform Experiments
- New Deep Learning (Engineers way)
 - Given huge amount of computational power
 - People First Experiment and then try to build a theory

Why Deep Learning ? Why Now ?

- **Computer Vision-** *Convolutional Neural Networks* and *Backpropagation* —well understood since 1989
 - **Time Series Forecasting-** *Long Short-Term Memory* — well understood since 1997
- “Deep Learning with Python” Francois Chollet

Why Deep Learning ? Why Now ?




Algorithmic Advancements...

- Better *Activation Functions* for neural layers.
- Better *Weight Initialization Schemes* starting with layer-wise pretraining.
- To avoid Overfitting the Concepts like *Dropout* is Introduced.
- Better *optimization schemes*, such as RMSProp and Adam.

Activation Functions...


- An *Activation Function (Transfer Function)* maps the weighted summation of inputs to output.
- An Activation function is used to add *Nonlinearity so that the network can learn complex patterns.*




Sigmoid Activation Functions

- **Characteristics:**
 - Differentiable
 - Nonlinear
 - O/P lies in [0-1]
 - Fast
 - *Vanishing Gradient Problem*

$$f(x) = \frac{1}{1 + e^{-x}}$$



$$f'(x) = g(x) = \text{sigmoid}(x) * (1 - \text{sigmoid}(x))$$



VANISHING GRADIENT PROBLEM


- Because of sigmoid activation function the derivative is *less than 1* and *when the derivatives are multiplied it gives a very small number* which ultimately changes the weight very less.
- *Usually occurs when the derivative is less than 1.*
- In case of *sigmoid and tanh activation* function it occurs frequently.

$$\frac{dL}{dw} = \frac{dL}{df_1} \times \frac{df_1}{df_2} \times \frac{df_2}{df_3} \times \dots \dots \dots \times \frac{df_n}{dw}$$

ReLU Activation Function


- $f(x) = x$, when $x > 0$
 $= 0$, when $x \leq 0$
- Avoids Vanishing Gradient Problem.
- Derivative is Simple
 - $f'(x) = 1$ for $x \geq 0$
 $= 0$ for $x < 0$
- Problem:
 - Dead ReLU Units

$$f(x) = \max(0, x)$$




Leaky ReLU Activation Function

- $f(x) = x$, when $x > 0$
 $= 0.1x$, when $x \leq 0$
- The advantages of Leaky ReLU are same as that of ReLU.
- In addition, it enables Backpropagation, even for negative input values.
- *Avoids Dead ReLU*
- Simple Derivative
 - $f'(x) = 1$ for $x \geq 0$
 $= 0.1$ for $x < 0$



WEIGHT INITIALIZATION




WEIGHT INITIALIZATION

- Mostly used
 - We should never initialize to same values.
 - Asymmetry is necessary
 - We should not initialize to large –ve values
 - Vanishing Gradient problems
 - Weights should be small (not too small)
 - Weights should have good variance
 - Weights should come from a Normal distribution with mean zero and small variance
 - Should have some +ve and Some –ve values

WEIGHT INITIALIZATION

- Better Strategies obtained from large experiments
 - Initialize weights based on Fan-in and Fan-out
 - Initialize your weights from a uniform distribution
 - $\left[-\frac{1}{\sqrt{fanin}}, \frac{1}{\sqrt{fanin}}\right]$
 - Works well for sigmoid activation function



WEIGHT INITIALIZATION

– Xavier/Glorot initialization in 2010- *well for sigmoid activation function*

- First Variation – $W_{ij} = N(0, \sigma_{ij})$, $\sigma_{ij} = \frac{2}{Fanin + Fanout}$
- Second Variation– $W_{ij} = U\left(-\frac{\sqrt{6}}{\sqrt{Fanin + Fanout}}, \frac{\sqrt{6}}{\sqrt{Fanin + Fanout}}\right)$

WEIGHT INITIALIZATION

– He Initializer, 2015 *works well for ReLU*

- First Variation – $W_{ij} = N(0, \sigma_{ij}), \quad \sigma_{ij} = \sqrt{\frac{2}{Fanin}}$
- Second Variation– $W_{ij} = U\left(-\frac{\sqrt{6}}{\sqrt{Fanin}}, \frac{\sqrt{6}}{\sqrt{Fanin}}\right)$


BIAS-VARIANCE TRADE-OFF



- Multilayer ANN has higher chance of overfitting.


DROPOUT AND REGULARIZATION

- Deep NN tend to overfit because of many layers and weights
- For this dropout and regularization is needed
- In Dropout, a certain percentage of inputs and hidden layer neurons are dropped out for an iteration
- Some call it as drop out network or layer.



Dropout

- Procedure:
 - During training we decide with probability p to update a node's weights or not.
 - We set p to be typically 0.5
- Highly effective in deep learning:
 - Decreases overfitting
 - Reduces training time
- Can be loosely interpreted as ensemble of networks



BATCH NORMALIZATION

- Normalization is a data pre-processing tool used to bring the numerical data to a common scale without distorting its shape.

- Decimal Scaling: $N_i = \frac{T_i}{10^p}$

- Median: $N_i = \frac{T_i}{\text{median}(T)}$


- Min-Max: $N_i = \text{Min}_N + \frac{T_i - \text{Min}_T}{\text{Max}_T - \text{Min}_T} \times (\text{Max}_N - \text{Min}_N)$

- Vector: $N_i = \frac{T_i}{\sqrt{\sum_{j=1}^k T_j^2}}$

- Z-Score: $N_i = \frac{T_i - \mu_T}{\sigma_T}$

BATCH NORMALIZATION

- Motivation



BATCH NORMALIZATION

$$\mu = \frac{1}{m} \sum h_i$$

$$\sigma = \sqrt{\frac{1}{m} \sum (h_i - \mu)^2}$$

$$h_{i(norm)} = \frac{h_i - \mu}{\sigma + \epsilon}$$


- Where m : Number of Neurons at h_i

$$h_i = \gamma \cdot h_{i(norm)} + \beta$$

- Where γ and β are hyper parameters.


BATCH NORMALIZATION

- Advantages
 - Faster Convergence
 - Weak Regularizer (Batch Normalization + dropout)
 - Avoids internal covariate shift
- <https://arxiv.org/pdf/1502.03167v3.pdf>






OPTIMIZERS

- At minima, maxima and saddle point, u have the gradient as Zero.




OPTIMIZERS

- Convex function and Non-Convex Function
- Convex functions have either 1 maxima or minima. (Local minima=global minima)
- Non-convex functions have more than one minima or maxima



Stochastic gradient descent (SGD)

You take one point (Input Vector), Feed Forward it then update the weights by back-propagating the gradient of errors.



- Initialize W_0 randomly
- For t in $0, \dots, T_{\text{maxiter}}$


$$W^{t+1} = W^t - \underbrace{\eta_t \cdot \nabla \text{Loss}(f_w(x_i), y_i)}_{\text{Stochastic gradient}}$$

where index i is chosen randomly

- computation of $\nabla \text{Loss}(\dots)$ requires only one training example
- Per-iteration comp. cost = $O(1)$

Gradient descent

You take all Input Vectors, Feed Forward it one by one, compute the error and get the mean error, then update the weights by back-propagating the gradient of errors.



- Initialize W_0 randomly

- For t in $0, \dots, T_{\text{maxiter}}$


$$W^{t+1} = W^t - \eta_t \cdot \underbrace{\nabla L(f_W(x_i), y_i)}_{\text{Gradient of the objective}}$$

Gradient of the objective

- computation of $\nabla L(W^t)$ requires a full sweep over the training data
- Per-iteration comp. cost = $O(n)$

Minibatch stochastic gradient descent

You take a subset of Input Vectors (more than one), Feed Forward it one by one, compute the error and get the mean error, then update the weights by back-propagating the gradient of errors.



- Initialize W_0 randomly

- For t in $0, \dots, T_{\text{maxiter}}$

$$W^{t+1} = W^t - \underbrace{\eta_t \cdot \tilde{\nabla}_B L(W)}_{\text{minibatch gradient}}$$

where minibatch B is chosen randomly

- $\tilde{\nabla} L(\theta)$ is average gradient over random subset of data of size B
- Per-iteration comp. cost = $O(B)$

STOCHASTIC GRADIENT WITH MOMENTUM

- The **rate of convergence** of Stochastic Gradient can be **improved** by *adding a momentum* to the Gradient expression.
- This can be *achieved by adding a fraction of previous weight change to the current weight change*.

$$(w_i)_{new} = (w_i)_{old} - \eta \left[\frac{dL}{dw_i} \right]$$

$$(w_i)_t = (w_i)_{t-1} - \alpha \cdot \Delta w_{t-1} - \eta \frac{dL}{dw}$$


Momentum

Learning
Rate

Nestrov Accelerated Gradient (NAG)


- SGD + Momentum

$$(w_i)_t = (w_i)_{t-1} - \alpha \cdot \Delta w_{t-1} - \eta \frac{dL}{dw}$$




Nestrov Accelerated Gradient (NAG)


- NAG



Nestrov Accelerated Gradient (NAG)



(a) Momentum-Based Gradient Descent



(b) Nesterov Accelerated Gradient Descent

ADAPTIVE GRADIENT(ADAGRAD)

- In SGD, SGD+Momentum and NAG, the learning rate is same for each weight.
- However, in Adagrad you have different learning rate for different weights.
- Why
 - Sparse Feature
 - Dense Feature

ADAPTIVE GRADIENT(ADAGRAD)

- SGD

$$(w_i)_{new} = (w_i)_{old} - \eta \left[\frac{dL}{dw_i} \right]$$


- Adagrad

$$\eta_t = \frac{\eta}{\sqrt{\alpha_{t-1} + \varepsilon}} \text{ with } \alpha_t \geq \alpha_{t-1}$$

$$\alpha_{t-1} = \sum_{i=1}^{t-1} \left(\frac{dL}{dw} \right)_i^2$$

- As iteration increases the learning rate decreases.

ADAPTIVE GRADIENT(ADA GRAD)



ADAPTIVE GRADIENT(ADAGRAD)

- Advantages
 - No need of manual tuning
 - Works well for both Sparse and Dense Feature
- Disadvantages
 - As iteration increases, the learning rate will get low, which will result in Slow Convergence.
 - Computationally expensive.

ADADelta

$$\eta'_t = \frac{\eta}{\sqrt{\text{Exponentially Decaying}(\alpha)_{t-1} + \epsilon}}$$

- $EDA_{t-1} = \gamma * EDA_{t-1} + (1 - \gamma) \left(\frac{dL}{dw} \right)_{t-2}^2$
- Avoids the Problem of slow convergence of AdaGrad

Root Mean Square Propagation (RMSProp)

- It is same to AdaDelta however, it discards the history from extreme past while computing the exponentially decaying average.
- Converges faster once it finds a locally convex bowl as its error function.
- Faster convergence than AdaDelta.

ADAM(ADAPTIVE MOMENT ESTIMATION)

- <https://arxiv.org/pdf/1412.6980.pdf>

$$w_{t+1} = w_t - \alpha m_t$$

where,

$$m_t = \beta m_{t-1} + (1 - \beta) \left[\frac{\partial L}{\partial w_t} \right]$$

m_t = aggregate of gradients at time t [current] (initially, $m_t = 0$)

m_{t-1} = aggregate of gradients at time $t-1$ [previous]

w_t = weights at time t

w_{t+1} = weights at time $t+1$

α_t = learning rate at time t

∂L = derivative of Loss Function

∂w_t = derivative of weights at time t

β = Moving average parameter (const, 0.9)

WHICH OPTIMIZER TO USE

- MiniBatch-SGD:::::::::: Small/Shallow ANN
- Momentum & NAG::: Works well in most cases but Slower
- AdaGrad:::::::::::::::::::: Sparse Features
- AdaDelta & RMSProp: Preferred Over AdaGrad
- Adam::::::::::::::::::::: Most Favorite

How to Train a Deep Neural Network?: Not Limited

- 1. Pre-processing: Data Normalization**
- 2. Weight Initialization**
 - Xavier & Glorot (For Sigmoid)
 - He Initializer (For ReLU)
- 3. Choose the Activation Function** (ReLU-Most Favourite)
- 4. Batch Normalization** (Especially for later layers close to O/P Layer)
- 5. Use Dropout**
- 6. Choose the Optimizer** (Favourite- Adam)
- 7. Hyper-parameters:** Architecture(# Layers, # Neurons), etc...
- 8. Loss Function**
 - 2-Class Classification : Log Loss
 - Multi-Class Classification: Multi-Class Log Loss
 - Regression: Squared Loss



For Your Valuable Time.